

A Multi-Purpose Recommender Framework

Halil I. Köklü

Supervised by Nigel Martin

April 2014

A Project Proposal Submitted to
Birkbeck College, University of London
in Partial Fulfillment of the Requirements
for the Degree of MSc Business Technologies

Department of Computer Science & Information Systems
Birkbeck College, University of London

Contents

List of Figures	3
1 Introduction	4
1.1 Proposal Structure	5
2 Background Research: Recommender Systems	6
2.1 Concepts	6
2.2 Techniques	7
2.2.1 Collaborative Filtering	8
2.2.2 Contentbased Filtering	9
2.2.3 Demographic Filtering	10
2.2.4 Knowledgebased Filtering	11
2.2.5 Hybrids	11
2.3 Adoption	12
2.4 Algorithms	13
2.4.1 Pearson Correlation Coefficient	13
2.4.2 Tanimoto Coefficient	14
3 Problem Definition	15
3.1 Framework Requirements	15
3.1.1 Multipurpose & Interoperability	15
3.1.2 Abstraction	16
3.1.3 Ease of Integration	19
3.2 Evaluation of Existing Research and Solutions	20
4 Technical Solution	22
4.1 Design & Architecture	22
4.1.1 Serviceoriented Architecture (SOA)	22
4.1.2 Multilayered Architecture	24

4.2	Technological Choices	27
4.2.1	Application Programming Interface (API)	27
4.2.2	Message Bus	29
4.2.3	Recommender Ecosystem	29
4.2.4	Others	30
4.3	Demo	31
5	Project Plan	33
5.1	Schedule	33
5.2	Fallback Plan	33
6	References	35

List of Figures

2.1	Collaborative Filtering	8
2.2	Item-to-Item Collaborative Filtering	9
2.3	Contentbased Filtering	9
2.4	Demographic Filtering	10
2.5	Knowledgebased Filtering	11
2.6	Feature Augmenting Hybrid Recommender System	12
2.7	Item-to-Item Recommendations by Amazon	13
2.8	Pearson Correlation Coefficient	14
2.9	Tanimoto Coefficient	14
3.1	Abstraction: Complexity	17
3.2	Abstraction: Missing Database Abstraction	17
3.3	Abstraction: Reusable Architecture	18
3.4	Abstraction: Dependency Issues	18
3.5	Abstraction: Encapsulation	19
3.6	Abstraction: Scalability	19
4.1	Service-Oriented Architecture (SOA)	22
4.2	Multilayered Architecture	24
4.3	Recommendation Model Configuration	27
4.4	Sample API for recommendations in <i>express</i>	28
4.5	Sample API message in <i>JavaScript Object Notation (JSON)</i>	28
4.6	Simple Graph in a Graph Database	30
4.7	Monitoring with NewRelic	31
4.8	Open-Source Ecommerce Web Application <i>Magento</i>	32
5.1	Project Schedule	34

1 Introduction

In this project I will design and develop an alternative architecture on the integration of recommender systems.

Essentially, a *recommender system* is a software which suggests useful items to a user on a given website or other platform. An item refers to any object available for the user and is subject to recommendations. A user is typically assisted in their search process on the platform which might be so vast that it becomes difficult to find the right item. Recommender systems are able to personalise so that it tries to suggest only items relevant for the particular user. In order to do this recommender systems rely on multiple techniques, notably *collaborative filtering* and *content-based filtering*.

Recommender systems are usually tightly coupled to the context and architecture of the beforementioned platform. Bespoke recommender systems may require to know and access the platform's database structure as well as fit into the same implementation constraints such as programming languages. This paper analyses issues arising from that problem and proposes an alternative approach to overcome these issues. This project aims to deliver a *multi-purpose recommender framework* which supports multiple techniques, is loosely coupled and easy to integrate. The framework provides an ecosystem with a plug-in system for recommender techniques and an event-driven learning concept. The solution of this project will be demonstrated in a real-life use case.

1.1 Proposal Structure

This proposal discusses specific challenges of recommender systems and – based on the background research as well as evaluation of existing solutions – proposes a potential solution to overcome them. Design requirements for the proposed solution are examined and resultant technology choices are defined. Finally, a project plan to ensure the project’s execution is compiled.

Section 2 introduces the concept of recommender systems. An overview of recommender techniques as well as algorithms are described.

Section 3 discusses challenges of recommender system and requirements of the proposed solution. Then, existing solutions are evaluated in their coping with these challenges.

Section 4 presents the architectural design of the solution and defines the technologies to be used.

Section 5 defines a project plan including methodology, evaluation, schedule and a fallback plan in case of complications.

2 Background Research:

Recommender Systems

This section introduces the fundamental concepts and techniques of recommender systems. Then, it discusses the adoption of recommender systems in research and enterprises. Finally, it gives two examples of algorithms commonly used in recommender systems.

2.1 Concepts

Recommendations are neither a new idea nor limited to the digital age. Ricci et al. (2011) write that traditional recommendations can be observed in various scenarios, such as a peer's recommendation when buying a book or reviews when choosing a movie. The authority of the recommender has an important role in the acceptance of the recommendation. A renowned film critic may appear more credible than a random colleague. When it comes to car parts, a mechanist may be a good candidate to ask. However the authority is not only limited to expertise – in fact we tend to rely more on recommendations which put our personal experiences and preferences into account. A previous companion for a wonderful trip can certainly have a better authority than a travel agent.

As mentioned in the introduction, with the growth of the Internet, the amount of information available on the Web increased rapidly. Especially, major e-commerce Web sites were extending their range of products and services. Although a wider and varied range of items is initially good for the user, users found it more and more difficult to find the appropriate items or make the right choices. Web sites

have deployed different type of solutions – such as search engines and more user friendly interfaces – to cope with this problem.

Another approach are recommender systems which basically provides a bespoke collection of items with the intention to highlight relevant items to a user. Depending on the recommending technique used various data sources are taken into account like the user’s context or previous interactions. This is a continuous learning process. The more the recommender system learns, the more accurate the recommendations will become. The user’s behaviour on given recommendations are further a powerful learning source for the system – e.g. if the user tends to accept some recommendations over others – to tweak the recommendations. Most recommender systems concentrate on guiding the user towards novel, unexperienced items (Herlocker et al., 2004).

2.2 Techniques

In the course of development different techniques to building recommendations emerged. Fundamentally techniques are classified by the information sources they use. The sources of personalised recommendations are typically user-item interactions (*collaborative filtering*), item features (*content-based filtering*), user features (*demographic filtering*) as well as knowledge about the user and item (*knowledge-based filtering*). Non-personalised recommendations also exist in form of ranked lists such as top sellers or related items. However they are typically not part of the research in recommender systems.

Anand and Mobasher (2003) differentiate between *explicit* and *implicit* data collection. Information the user intentionally provides to the system to express a positive preference to items are referred to as *explicit* data collection including rating an item, adding it to a wish list or liking it. *Implicit* data on the other hand is collected by observing the user’s behaviour e.g. usage of navigation and search elements or purchase of items.

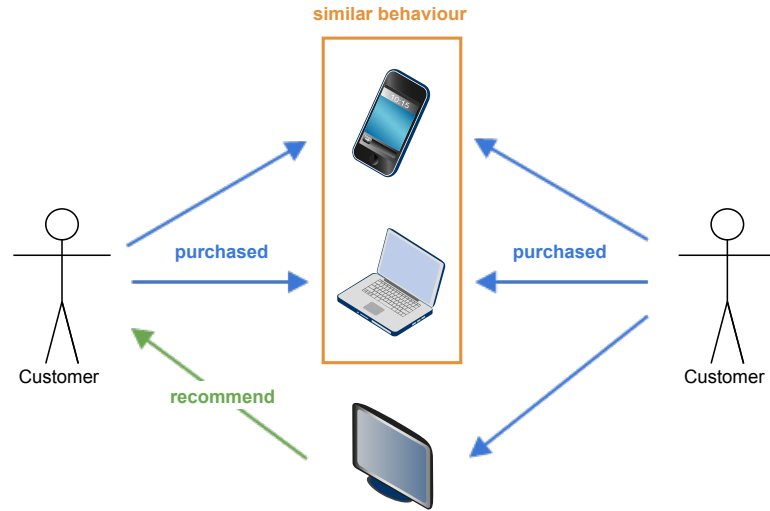


Figure 2.1: Collaborative Filtering

2.2.1 Collaborative Filtering

Basically this technique recommends items other users with similar preferences interacted with e.g. liked or purchased. In order to do this the recommender system needs to observe users' behaviours and interactions. Based on these learnings, it will then look up other users with similar behaviour patterns and build recommendations from their preferences – preferably items which the active user has not experienced yet (see figure 2.1).

The major advantage of the collaborative filtering approach is that the recommender system does not require any knowledge about the items.

Figure 2.1 illustrates a scenario where the active customer has purchased several items in the past. The recommender system understands that the active customer is similar to another as both have purchased the phone and laptop. Then, the system computes items the similar customer purchased but the active customer has not. It therefore recommends the TV to the first customer.

Item-to-Item Collaborative Filtering

This approach is a derivative of traditional collaborative filtering methods and has been popularised by Amazon (Linden et al., 2003). The fundamental difference lies in the fact that the item-to-item approach collects collaborative data to put

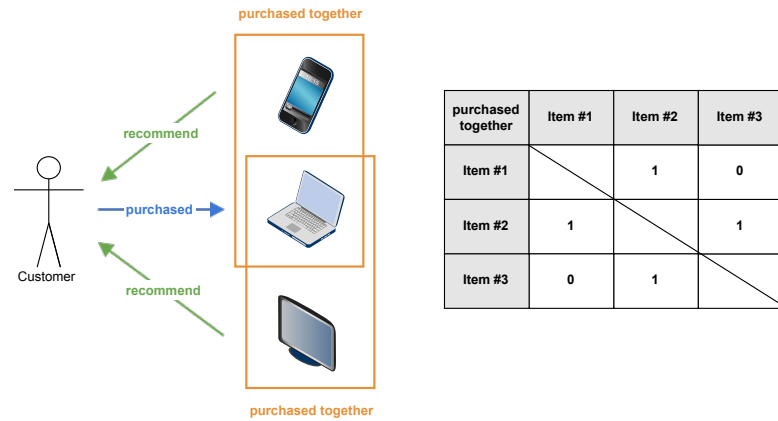


Figure 2.2: Item-to-Item Collaborative Filtering

items into relation with other items rather than users. The recommendation query takes an item – currently viewed or from a wish list – as an argument and looks for items purchased together with that item (see figure 2.7).

Given that the item range can be very wide, item-to-item filtering methods require significant computing time and data storage. However they are usually preprocessed offline thus queries can be processed rather quickly.

Figure 2.2 demonstrates a customer who has purchased an item in the past. The recommender system looks up items which were purchased together with that item and recommends them to the customer.

2.2.2 Contentbased Filtering

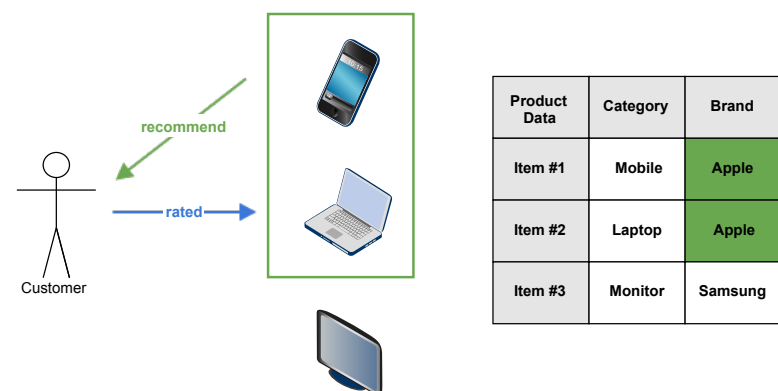


Figure 2.3: Contentbased Filtering

Content-based recommendation methods make use of item features to find similar items. Based on items the user has showed a positive preference to before –

such as rated or purchased – other similar items are looked up based on the item’s features. Figure 2.3 illustrates a customer who has rated an item positively. The recommender system compares the rated item with other items, finds another item which has the same brand and therefore recommends that item.

2.2.3 Demographic Filtering

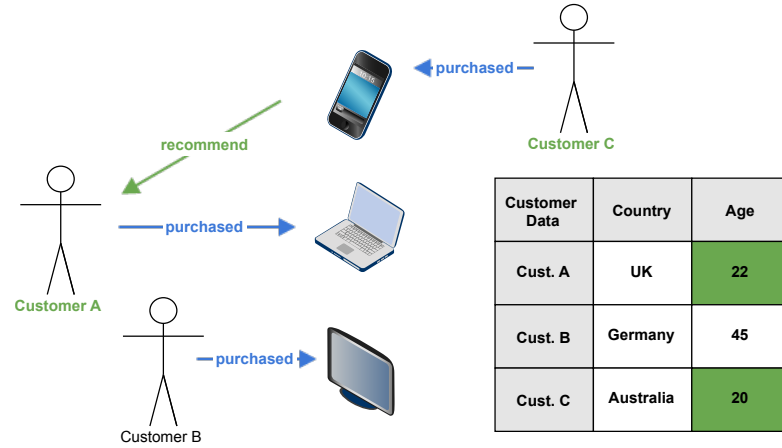


Figure 2.4: Demographic Filtering

The demographic filtering is similar to the content-based method with the significant difference that demographic filtering examines the demographic profile of a user – user features rather than item features – to find similar users and recommend items those users have showed positive preference to in the past. Burke (2007) makes the assumption that recommendations should be different for demographic groups. The demographic profile can consist of age, gender, interests, language, country etc. To give an example, a hotel search engine may want to recommend different hotels to a business person and different ones for a young couple.

Figure 2.4 shows how a demographic filtering recommender system would recommend an item to customer A based on a customer base of three customers with previous purchase history. The recommender analyses the customer profile features and eventually decides relying on the age feature that customer C is similar to customer A. Hence the system recommends a product customer C purchased in the past.

2.2.4 Knowledgebased Filtering




Figure 2.5: Knowledgebased Filtering

Recommendations of knowledge-based systems rely on specific domain knowledge to determine useful items for a user. These systems are usually constraint-based or case-based. Both approaches are similar in their conversational process. In other words, the user specifies the requirements and the system tries to find solutions which are fulfilling them. Whereas constraint-based systems are matching the requirements explicitly (such as price ranges), case-based systems make use of similarity measures (e.g. distance to a point of interest) (Ricci et al., 2011).

An example of a knowledge-based recommender is figure 2.5 which recommends a monitor as a peripheral to a customer who views a laptop. This is based on explicit knowledge that amongst others external screens are a demand for laptop users.

2.2.5 Hybrids

Hybrid recommender systems make use of two or more individual recommender systems – hereinafter referred to as components – to combine aforementioned techniques. One possible motivation of using hybrid systems may be to overcome weaknesses of one approach by combining it with another. However it is also possible to combine systems implementing the same technique but e.g. using different knowledge sources. There are many possible ways of combining techniques of which Burke (2007) identified seven:

Weighted The scores of recommender components are combined using a linear formula. A score is a numerical rank attached to items.

Switching The system chooses one component over another which is based on some criterion.

Mixed Recommendations from components are summed up and presented together. Implementations usually differ on how the results are merged together.

Feature Combination A contributing component modifies the features of a knowledge source and feeds into the actual recommender component.

Feature Augmentation Similar to *feature combination* this approach's contributing component adds new features rather than modify them (see figure 2.6).

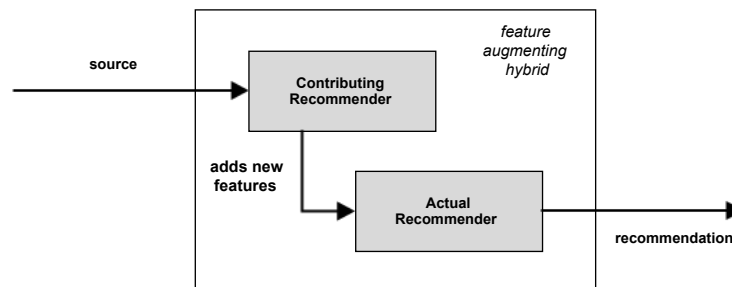


Figure 2.6: Feature Augmenting Hybrid Recommender System

Cascade Components are given strict priorities with the lower prioritised ones breaking ties of the higher ones.

Meta-Level Similar to *feature combination* as well as *feature augmentation*, yet the contributing component completely replaces – instead of appends or modifies – the initial source.

2.3 Adoption

Ricci et al. (2011) note that research on recommender systems is relatively new compared to other classical information retrieval methods like databases and search engines. However it is gaining attention due to various reasons. First of all e-commerce companies including Amazon, YouTube and Netflix invest a lot

in recommender systems. As trendsetters and pioneers for other institutions they build a demand for further research and development. Dedicated conferences and academic courses as well as its adaption in several academic journals are indicators of recognition of this research area.

E-commerce companies running recommender systems hope among others to increase number of items sold, sell more diverse items, increase customer satisfaction and recommend sequences plus bundles (Herlocker et al., 2004).

Inspired by Your Wish List

You wished for	Customers who viewed this also viewed
 <p>The Innovator's Solution: Creating... ▶ Clayton M. Christensen, Michael E. Raynor Hardcover ★★★★★ (89)</p>	 <p>The Effective Executive: The... ▶ Peter F. Drucker Paperback ★★★★★ (152) \$46.99 \$9.60</p>
	 <p>The Innovator's Solution: Creating... Clayton M. Christensen, Michael E... Hardcover ★★★★★ (7) \$30.00 \$20.73</p>
	 <p>The Innovator's Dilemma: The... Clayton M. Christensen Paperback ★★★★★ (133) \$47.99 \$10.16</p>

Figure 2.7: Item-to-Item Recommendations by Amazon

2.4 Algorithms

Recommender systems rely on algorithms. They are usually defined to solve a specific problem and it is therefore crucial to utilise the right algorithm to get the expected results.

In the next two sections I will illustrate two algorithms which are very popular in recommender systems.

2.4.1 Pearson Correlation Coefficient

The *Pearson correlation coefficient* provides a measure to calculate the correlation of two variables. In the recommender systems context it is useful to compare user's behaviour patterns or preferences (Segaran, 2007). The coefficient is defined to be

$$Pearson(a, b) = \frac{\sum r_a r_b - \frac{\sum r_a \sum r_b}{N_{ab}}}{\sqrt{(\sum r_a^2 - \frac{(\sum r_a)^2}{N_{ab}})(\sum r_b^2 - \frac{(\sum r_b)^2}{N_{ab}})}}$$

Figure 2.8: Pearson Correlation Coefficient between ratings by the users a and b where r_a respectively r_b are rating vectors for the mutually rated items N_{ab} .

between 1 to -1 where 1 indicates a perfect correlation, 0 no correlation and -1 a perfectly inverse correlation.

2.4.2 Tanimoto Coefficient

$$T = \frac{N_{ab}}{(N_a + N_b - N_{ab})}$$

Figure 2.9: Tanimoto Coefficient where N_a respectively N_b is the count of properties of a respectively b and N_{ab} is the count of intersecting properties.

The *Tanimoto coefficient* tells us the similarity of two sets (Segaran, 2007). It can be used to measure how similar two items or users are based on their features which is useful for *content-based* and *demographic filtering* methods. Below is an example of two items which we want to compare:

```
A = [mobile, apple, iphone, black, 32G]
B = [mobile, apple, ipad, black]
```

Given the equation in figure 2.9 we come to the conclusion:

$$T = \frac{N_{ab}}{(N_a + N_b - N_{ab})} = \frac{3}{5 + 4 - 3} = 0.5$$

3 Problem Definition

The primary objective of this project is to implement a *multi-purpose recommender framework* to overcome specific issues of recommender systems explained further in this section. The focus on *framework* intends to emphasise an ecosystem which provides a distinctive tool kit to build multiple recommender systems, yet guarantees a simple integration. Insofar the focus of the project will be in the design of an alternative – perhaps novel – architecture rather than in the details of techniques and algorithms.

This section discusses the requirements of the proposed framework while pointing to aforementioned issues of recommender systems. The requirements *interoperability*, *abstraction*, *ease of integration* and *service-orientation* will become measures on how successful the proposed solution is in overcoming those issues.

Finally, existing research as well as solutions are examined in their coping with those challenges.

3.1 Framework Requirements

3.1.1 Multipurpose & Interoperability

A *multi-purpose* recommender system is able to cope with different data sources and techniques in parallel. Ideally this system is open and extendable for possible forthcoming, yet unknown requirements. It is not to be confused with hybrid recommender systems (see section 2.2.5) which combine different systems into one. In this project I will provide an ecosystem to plug different recommender systems in (see section 4.1.2).

Another requirement is the interoperability between recommender systems. Manouselis and Costopoulou (2007) differentiate between three criteria:

Interoperability of the recommendation queries which allows the same query to be reused. This is in particular useful for hybrid recommender systems which feeds the same query into different systems. Adomavicius et al. (2005) go further and creates a *Recommendation Query Language (RQL)*. The use of an RQL would go beyond the scope of this project. Therefore I will limit this criterion to query parameters rather than the full query and rely on the recommender system to build the query based on these parameters. To give an example, the query parameters for figure 2.3 would be the customer reference and the rated item (in this case the laptop). A content-based recommender system would compute these parameters by looking up similar items whereas a collaborative filtering system would look for other users who rated this item.

Interoperability of the user and the domain models which allows the exchange of models and other data among different recommender systems. In this project I will go a step further and allow direct access to persistence layers to all recommender systems.

Interoperability of the recommendation results which empowers recommender systems to reuse the results. This is in particular useful for hybrid recommender systems which feeds the same query into different systems. Yet there is no use case in this project and therefore not covered.

3.1.2 Abstraction

Recommender systems are usually tightly coupled to the primary application, which is the main application using the recommender system. Cortizo et al. (2010) claim that recommender systems are designed for a specific data structure. This comes with the following challenges:

Complexity Applications as well as recommender systems tend to be complex. If not separated they both add to the overall complexity which makes change

management very difficult, time consuming and expensive. Any change requires knowledge – higher human resources costs –, testing and probably modifications on both systems. Figure 3.1 shows a tightly coupled system whose one recommender system is getting replaced. The whole primary application as well as the database are getting affected.

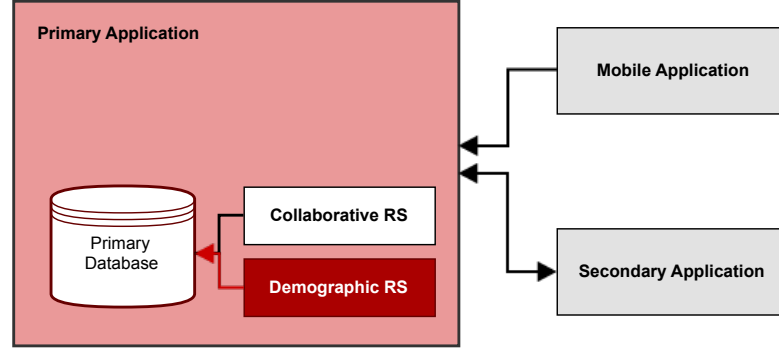


Figure 3.1: Abstraction: Complexity

Database Abstraction Recommender systems directly using databases of the primary or other applications are problematic as on one hand they are affected to any schema changes done to the database due to application modifications. Figure 3.2 illustrates a semi-decoupled architecture where the recommender system is outside of the primary application, yet still uses its database. This gives a false impression of loose coupling. Teams working on the primary application might not fully aware that there changes affect other systems.

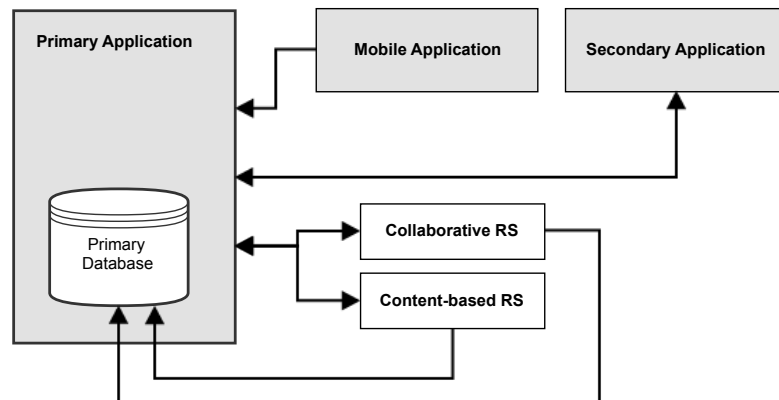


Figure 3.2: Abstraction: Missing Database Abstraction

Reusability Tightly coupled components are difficult to reuse. Figure 3.3 shows an architecture where other applications and even devices such as mobile

can reuse the recommender systems.

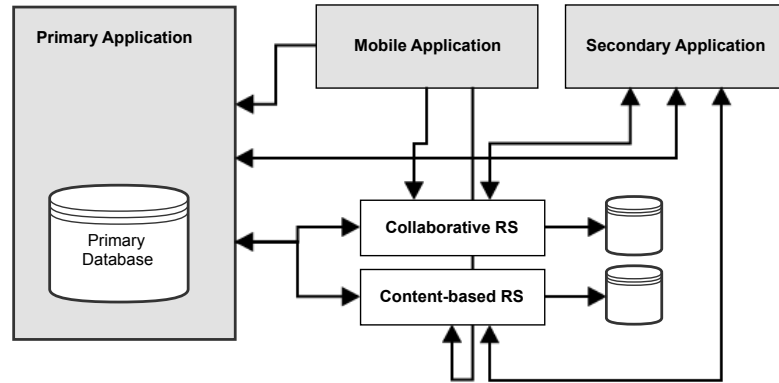


Figure 3.3: Abstraction: Reusable Architecture

Dependency If a system needs to be replaced or modified, it can have more or less major implications to other systems depending on that system. Figure 3.4 shows a loosely coupled architecture where the recommender technique has been changed from content-based to demographic filtering. Due to the loose coupling it has less implications than in figure 3.1. However it can still affect other systems if the communication between them included content-based filtering specific logic.

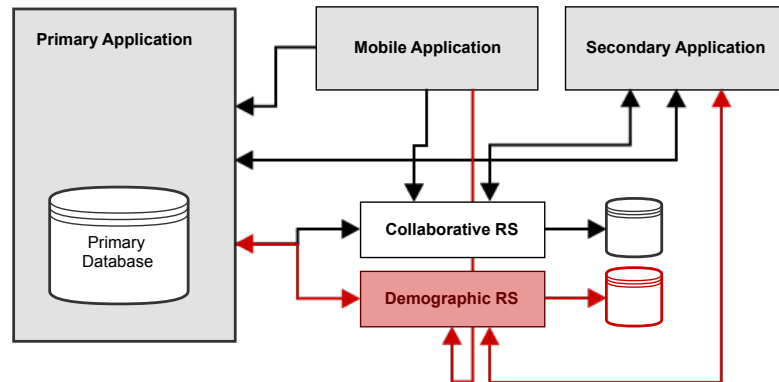


Figure 3.4: Abstraction: Dependency Issues

Encapsulation Information hiding is the fundamental motivation for encapsulation. The more information and implementation is hidden, the looser the coupling becomes. Figure 3.5 illustrates a recommender system framework which hides internal details and communicates with other components in an implementation unspecific manner. A change of a technique within the framework or an external application should not have any implications.

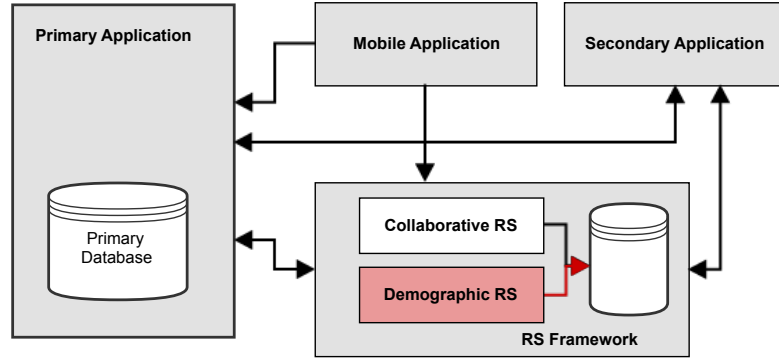


Figure 3.5: Abstraction: Encapsulation

Scalability Different components have different problems regarding scalability. In terms of scalability two key scaling approaches exist. *Horizontal scaling* suggests adding more nodes such as servers to the system, whereas *vertical scaling* adds more resources such as memory to an existing system. Tightly coupled systems are difficult to scale. E.g. if the primary application has high *random-access memory (RAM)* usage and the recommender system requires higher *central processing units (CPU)*, it is not directly possible to scale horizontally by adding nodes with sufficient CPU. They also need to have enough RAM which is usually more expensive. On another note, the whole system may not support horizontal scaling at all due to design problems in the complexity.

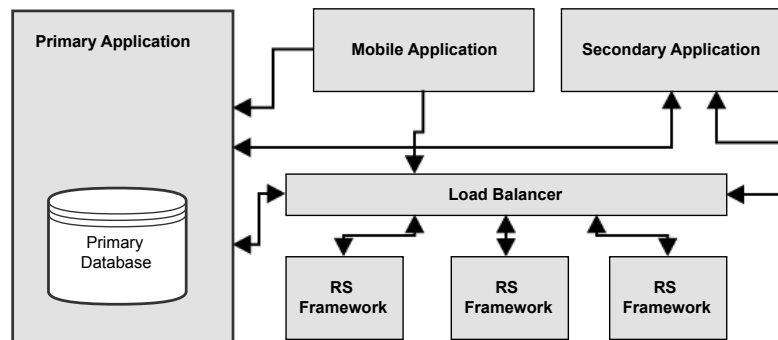


Figure 3.6: Abstraction: Scalability

3.1.3 Ease of Integration

The complexity and cost of integrating recommender systems is a major constraint for many projects. Given the requirements above it is even more challenging to

keep the integration as easy as possible. Nonetheless abstraction should help reduce the amount of knowledge needed. Furthermore the fewer types of communication and configuration required the less effort will be necessary.

3.2 Evaluation of Existing Research and Solutions

As mentioned in 2.3, recommender systems are considered a relatively new research area in computer science. Although it has gained attention recently, a considerable amount of research has been published on the fundamental concepts and techniques rather than on matters of integration and architecture. Nonetheless I found that Cortizo et al. (2010) and Rack et al. (2007) worked on similar approaches.

Cortizo et al. (2010) built a *general purpose multi-algorithm* recommender system to compute recommendations from different sources and serving multiple applications. They mention that they could not find any literature on the system's aspects of recommender systems. Cortizo et al. (2010) identified similar challenges and made similar decisions on the implementation especially using RESTful APIs (see section 4.1.2). The advantage of their work was that they evaluated and tested their system on a live environment. Scalability and performance were key metrics from the beginning. Their recommender system is not open source and therefore not available for me.

Rack et al. (2007) have published a series of papers around their work on the *AMAYA* recommender system. Although they have put an emphasis on multipurposeness, their work is more oriented towards a context-aware recommender system which differentiates between contexts such as '*being home*' and '*being at work*'. Furthermore their recommender requires a user profile as a centre point. In summary, their architecture is not as flexible and unbiased as the design I propose. The last submitted paper about *AMAYA* was in 2007 and the system is again not publicly available.

There are several commercial recommender software available such as *prudsys*

Realtime Decisioning Engine (RDE) which I have integrated into a major e-commerce website in the past. In a brief analysis of commercial software I found that they usually cover only basic but common requirements for e-commerce websites. These products are closed source and therefore not beneficial for my project.

Hahsler (2011) and Rack et al. (2007) provide a list of open source recommender systems which are freely available. In the majority these systems are more component libraries than complete solutions. They require significant amount of work to integrate. Amongst them is *Apache Mahout* – a machine learning library which also includes collaborative filtering components. *Mahout* is built upon *Hadoop* which is a well-established big data software also curated by *The Apache Software Foundation*. *easyrec* on the other hand is a complete solution which is also using RESTful APIs. However – similar to commercial products – it is opiated in favor of e-commerce websites.

4 Technical Solution

In this section an overview of the design, architecture and distinctive features of the solution is given. Then, technical decisions are explained.

4.1 Design & Architecture

The architecture of this project is designed to meet the requirements discussed in section 3. First, a bird's eye view is given on where this recommender framework fits into existing systems. Then, the internal architecture of the framework is explained. Finally, distinctive features of this framework are introduced.

4.1.1 Serviceoriented Architecture (SOA)

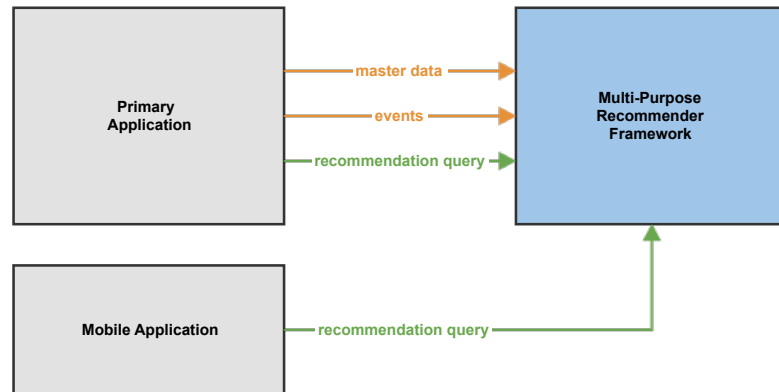


Figure 4.1: Service-Oriented Architecture (SOA). Orange arrows (*master data* and *event*) are *notifications* whereas green arrows (*recommendation query*) are *queries* and expect a result.

A *service-oriented architecture (SOA)* is a software design pattern which is most suitable to meet the *abstraction* requirement discussed in section 3.1.2. *SOA* sug-

gests to express features as services. This is true for features which are going to be available to other systems. Internal features are never to be exposed and allowed to be services. A *service consumer* is a system using the service. (Erl, 2008) identifies eight principles of *SOA* of which I will elaborate:

Standardised service contracts is an expression of the service's purpose, capabilities and requirements – such as mandatory parameters and data types. As long as the requirements are satisfied, the service agrees to fulfill its purpose.

Service loose coupling makes sure that services have as few dependancies as possible.

Service abstraction ensure that as much information as possible hidden and none except those described in the service contract are exposed.

Service reusability assure that services are designed to be reused.

In figure 4.1 a possible architecture of this project is shown. The framework is designed to serve more than one application or system component. The services accept requests from any source as long as they authorise with an access key. In the mentioned figure services are divided in two types – *notification* and *query*. A *notification* is a message relevant to the recommender framework and is simply acknowledged as received. A *query* on the other hand expects a result. The framework will define three services:

Master Data enables service consumers to create, update or delete a data node in the framework. An identifier and type are mandatory fields. The service consumer is allowed to send any further features of the node which it thinks is relevant to the framework.

Event accepts notifications about interactions or preferences between two or more nodes such as '*X purchased Y and Z*'. The payload – content of the message – may contain a *weight* which is a numeric value. It is useful for e.g. '*X rated Y with 10*'.

Recommendation Query requests recommendations for a specific *recommendation model*. The recommendation model identifier is mandatory. A node identifier and type are only mandatory if the model expects it. This is the

only service which expects a result.

The recommendation framework only expects incoming requests (*push strategy*) and has no outgoing communication at all.

4.1.2 Multilayered Architecture

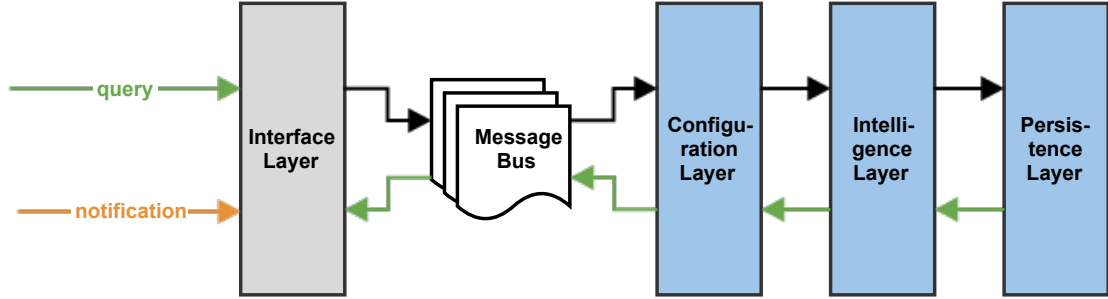


Figure 4.2: Multilayered Architecture

A multi-layered architecture is a software design pattern for software programs which suggests separating functionality in responsibility layers. This is popular way of abstracting functionality within a system.

Strictly speaking, the framework is divided into two subsystems: the *application programming interface (API)* and *recommender ecosystem*. These subsystems are connected with a message bus. The motivation of this differentiation is of a technical rather than logical manner. The subsystems have different kind of technical requirements and this division allows me to use the most appropriate technology for each. Another side effect are the advantages of a message bus which allows me to group and prioritise messages. *Recommendation queries* may be processed by a dedicated resource whereas *event* messages may have higher priority – and should be processed sooner – than *master data* messages. However this makes no difference to the responsibilities, and is therefore not further considered in the specification of layers.

Interface Layer

The interface layer is where services expose themselves to service consumers. Latter send requests to this layer. From a technical point of view this layer is imple-

mented as an *application programming interface (API)* via the *hypertext transfer protocol (HTTP)* which is the foundation of data communication in the world wide web. The *API* adopts the *representational state transfer (REST)* style. An *API* using *REST* is called *RESTful*. Fielding (2000) – who introduced *REST* – explains:

The name “Representational State Transfer” is intended to evoke an image of how a well-designed Web application behaves: a network of web pages (a virtual state-machine), where the user progresses through the application by selecting links (state transitions), resulting in the next page (representing the next state of the application) being transferred to the user and rendered for their use.

Two concepts of *REST* are important to understand: the use of *uniform resource locator (URL)* and *HTTP vocabulary*. The *URL* provides a way to specify a resource – in this case a service or node. The *HTTP vocabulary* defines amongst others *GET*, *POST* and *DELETE*. Combined a *RESTful API* enables the access or modification of resources. E.g. to submit an event a *POST* request to */events* is necessary. To delete a node with the identifier *120* a *DELETE* request to */nodes/120* is sufficient, whereas a *GET* request to */recommendations/topseller* would return all recommendations for a recommendation model called *topseller*. By using low-level *HTTP* vocabulary the *API* requires less documentation and explanation.

This layer acts as *service broker* and *security agent* at the same time. Latter is ensured by verifying the presence and correctness of an access key which is a random text. As a *service broker* this layer validates the payload against mandatory fields and data types. Then, it offloads the message and submits it into the message bus for further processing. If the message is of the type *notification* it acknowledges the request. If the message type is a *query*, then it waits for a response from the message bus and returns that.

Configuration Layer

This layer is the heart of the multi-purpose recommender framework and targets to satisfy the multi-purpose and interoperability requirement (as defined in 3.1.1). The motivation behind this layer is to empower operators – who use this solution – to set up a recommender system by configuration only. The configuration layer consists of two fundamental ideas: the *recommendation model* and *event* subsystems.

A *recommendation model* is an instruction to the recommender to compute and how to compute recommendations. In the model configuration the recommender to be used is specified. An identifier allows service consumers to access recommendations for that particular model.

As mentioned in section 4.1.1, an event is an interaction or preference between two or more nodes such as '*X purchased Y and Z*'. The event subsystem provides a mechanism to allow recommendation models to listen to events. When an event message is received, the subsystem will update listeners about it. This way recommendation models receive new data. One event can be interesting for more than one recommendation model. In a traditional architecture the primary application would need to update every single recommender system. However the event subsystem makes redundant API calls obsolete as only one message is sufficient, to update many recommendation models. Figure 4.3 shows an example configuration.

Recommendation Layer

In the *recommendation layer* the actual recommender algorithms will remain. These algorithms are called explicitly based on the recommender specified by the configuration layer. The recommender algorithms are going to be unaware of the nature of the data. In the contrary it will work based on the parameters fed by the configuration layer. Then, the algorithm will build its algorithm-specific query and use the persistence layer to fetch data and eventually return recommendations.

A distinctive feature of this framework will be the recommender plug-in sub-

```

<config>
  <rule name="items_bought_together" recommender="item_amazon">
    <on event="order_created" do="update" />
    <on event="product_removed" do="remove" />
  </rule>
  <rule name="similar_customers" recommender="collaborative">
    <on event="productRated" do="add" />
    <on event="order_created" do="add" />
  </rule>
  <rule name="similar_products" node_type="product" recommender="content_based" />
</config>

```

Figure 4.3: Recommendation Model Configuration

system. All recommenders will be designed as extensions. This subsystem will support extending the framework with new recommenders. This is actually a big potential to become a mechanism to define hybrid recommenders based on other recommenders in the framework and define a combination criterion (as discussed in section 2.2.5). However most probably time would not permit to work on the latter.

Persistence Layer

This layer provides database abstraction components which will then allow any recommender system to access any of the supported storage systems.

4.2 Technological Choices

In this section I will discuss my choices in technology I want to use in this project.

4.2.1 Application Programming Interface (API)

The *API* is the interface layer in this project and has almost none business logic. In that sense it is important to use a lightweight, thin solution. As discussed in

4.1.2 I will use a *RESTful* approach for the *API*.

The *API* will be built on the *node.js* platform which makes use of *Google Chrome*'s fast *V8 JavaScript engine* and requires only a few lines of code to create a server. The *API* will further use *express* – a web application framework for *node.js*. *express* is ideal for *RESTful APIs* as it uses the *HTTP* vocabulary as well. Figure 4.4 illustrates a sample implementation of ‘*GET* request to */recommendations/topseller*’ mentioned in 4.1.2. As visible in the figure the footprint of the implementation is very thin. As the *API* has a specific format in *express*, it is able to automatically generate *API* documentation.

```
app.get('/recommendations/:id', function(req, res){  
    // do work  
});
```

Figure 4.4: Sample *API* for recommendations in *express*

The message format will be in *JavaScript Object Notation (JSON)* - a thinner alternative to *extensible markup language (XML)* (listing 4.5).

```
{  
  "nodes": [  
    {  
      "id": 10,  
      "type": "product",  
      "title": "Laptop"  
    },  
    {  
      "id": 20,  
      "type": "user",  
      "name": "halil"  
    }  
  ]  
}
```

Figure 4.5: Sample *API* message in *JavaScript Object Notation (JSON)*

4.2.2 Message Bus

The concept of the message bus was introduced in section 4.1.2. The main requirement for this message bus is that it implements the *advanced message queuing protocol (AMQP)* – an open standard which defines a minimum set of features in message buses. A popular, highly reliable message bus software which supports *AMQP* is *RabbitMQ*. It also provides a *graphical user interface (GUI)* giving information about the current message flow. I have worked with *RabbitMQ* in the past.

4.2.3 Recommender Ecosystem

This subsystem contains configuration, intelligence and persistence layers. It has the highest technical requirements amongst the solution as it needs to be performant and stable. Especially, the recommender algorithms will require a performant environment. In that sense dynamic scripting languages such as *hypertext preprocessor (PHP)* or *Python* are not my first choice. In fact a simple statically-typed, compiled language with concurrency capabilities is needed.

My preferred candidate for this is *Go* – a language developed by *Google* as an alternative to overcome limitations of the programming language *C++*. The result is a language which is intendedly not adopting all patterns found in many programming languages such as overloading and pointers. The language aims to be simple, safe and free of misinterpretations by the compiler. It supports multithreading, CPU paralleling as well as asynchrony. Finally, it has a modern package management system which allows retrieving packages from the internet.

Database Software

The recommender ecosystem requires a powerful yet flexible database software. As pointed out in section 4.1.1 a master data can have an arbitrary number of additional fields. To fulfill the *ease of integration* requirement a schema management solution is not desired. Schema-less databases – also known as *NoSQL* –

might be of interest. *NoSQL* databases are usually non-relational as well. However recommendations are very intensive in terms of relations (e.g. ‘*X rated Y*’).

A potential solution are graph databases which understands *nodes*, *properties* and *edges*. A *node* is a reference with an identifier. An *edge* is a path – thus relation – from one *node* to another. A *property* can be attached to *nodes* as well as *edges*. Latter is useful for weighted relations such as ‘*X rated Y with 10*’. Figure 4.6 shows a sample graph. Graph databases have further advantages especially in querying distant *nodes* via other *nodes*. With regard to figure 4.6, an example query could be to fetch all groups people, who Alice knows, are members of.

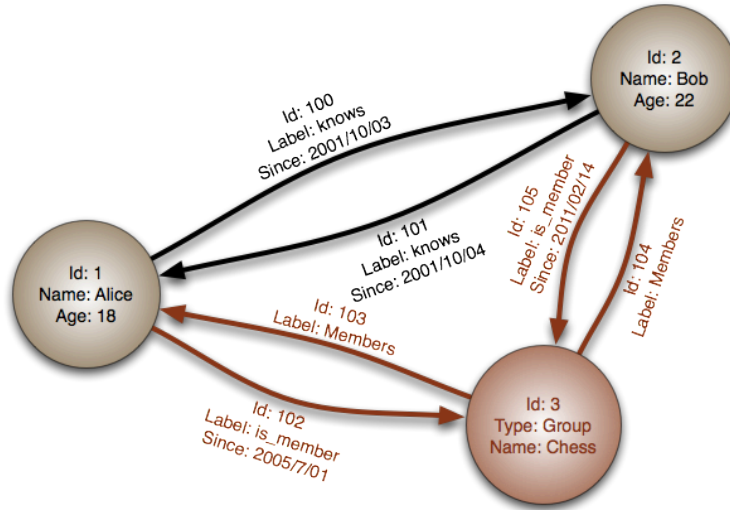


Figure 4.6: Simple Graph in a Graph Database. Source: Creative Commons.

If time permits an additional database system – called *Redis* – is evaluated. *Redis* is a *key-value* storage where a *value* is stored with a reference called *key*. There are no relations. A typical query of an *item-to-item collaborative filtering* method (see section 2.2.1) would be ‘*fetch all products bought together with X*’. The recommender framework would therefore make two storage solutions available.

4.2.4 Others

A *version control management (VCS)* namely *git* will be used throughout the project. *Git* is a distributed *VCS* which is amongst others faster and more flexible than other *VCS* such as *Subversion*. The *git* repository – and therefore the source

code – will be hosted on *BitBucket* to have a backup of the project anytime.

For server health and performance monitoring I will use the *software as a service* (SAAS) solution *NewRelic*. It also allows to analyse specific low performing or faulty requests. This will be very useful during testing and evaluation of the solution.




Figure 4.7: Monitoring with NewRelic

Finally, a virtualisation solution called *VirtualBox* will be used to set up all technical and vendor requirements within a virtual machine. This prevents conflicts with dependencies on the workstation especially if versions differ. When submitting the project the virtual machine will be packaged to provide a seamless demo setup for examination.

4.3 Demo

The proposed architecture primarily deals with internals. Although it can be tested by simulating events, it is not very visible and tangible. Therefore the *multi-purpose recommender framework* will be demonstrated on a platform which reflects real-life use cases. I will try to showcase as many different techniques – such as collaborative and content-based filtering. As mentioned in the background research most applications using recommender systems are e-commerce websites. It is hence not surprising that one of the demos is covering that area.

The storefront will be based on the open source e-commerce web application *Magento* – strictly speaking the free community edition. Having used *Magento* professionally I am comfortable extending it for my testing purposes. The product database will be loaded with ca. hundred thousand individual products which allows us scalability and performance testing of the framework APIs as well.

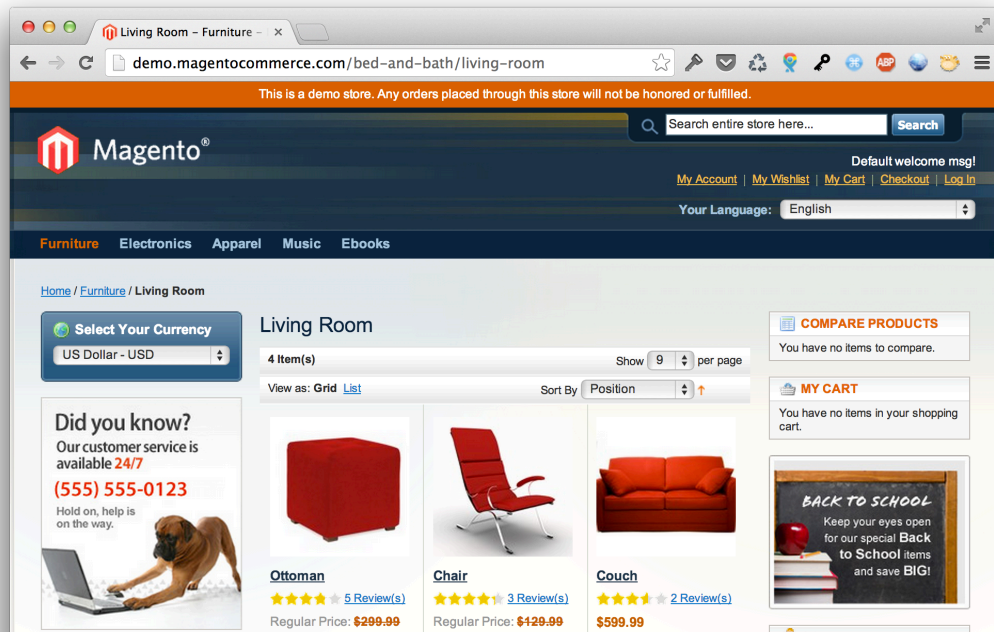


Figure 4.8: Open-Source Ecommerce Web Application *Magento*

5 Project Plan

This section provides a work schedule to complete this project as well as a fallback plan in case of unexpected or time-related circumstances.

5.1 Schedule

This project will be split into two broad phases: *implementation* and *evaluation*. The implementation phase is from 1st May to 31st July 2014, then the evaluation phase is until 15th September 2014. As the project has significant implementation work to be done, this strict division makes sure that enough time is left for the evaluation phase.

The schedule is divided into bi-weekly sprints with the aim that at the end of each sprint, a completed part of work is delivered. For the given time frame, there are nine sprints. Figure 5.1 shows the expected deliverables in these sprints.

5.2 Fallback Plan

In this section alternative paths are defined in case of delays due to underestimated work or other external circumstances.

The software architecture is the critical path of this project. Little is gained when layers or distinctive features are missing. Therefore I will focus on delivering them. In case of delays it is possible to decrease the number of different techniques and recommendation scenarios. In the worst case scenario the abstraction level of the recommendation model can be lowered so that recommender techniques are not as generic as purposed.

	Sprint	Deliverables
Implementation	1	Set up virtual machine as well as Magento. Implement interface layer. Implement publishing to message bus.
	2	Implement master data service.
	3	Implement event service.
	4	Implement plug-in system for recommender techniques.
	5	Implement recommendation model service.
	6	Implement several recommender techniques
Evaluation	7	Evaluation of architecture and recommendation quality. Structure report.
	8	Write implementation related chapters of report.
	9	Write critical evaluation related chapters of report. Write documentation.

Figure 5.1: Project Schedule

Go and *Neo4j* are unfamiliar to me. In case of slow progress due to the learning process, they maybe replaced by simpler or more familiar choices. *Go* could be replaced by *PHP* or *Python*. The usage of *Redis* is planned only if time permits. However if problems arise with *Neo4j*, *Redis* may be used as alternative.

6 References

- Adomavicius, G., Tuzhilin, A., and Zheng, R. (2005). Rql: A query language for recommender systems.
- Advanced Message Queuing Protocol (AMQP) (2014). <http://www.amqp.org/>. Accessed April 07, 2014.
- Amazon.com (2014). <http://www.amazon.com/>. Accessed April 07, 2014.
- Anand, S. S. and Mobasher, B. (2003). Intelligent techniques for web personalization. *Proceedings of the 2003 international conference on Intelligent Techniques for Web Personalization*, 7(4).
- Apache Hadoop (2005-2014). <http://hadoop.apache.org/>. Accessed April 07, 2014.
- Apache Mahout (2009-2014). <https://mahout.apache.org/>. Accessed April 07, 2014.
- BitBucket (2014). <https://bitbucket.org/>. Accessed April 07, 2014.
- Burke, R. (2007). Hybrid web recommender systems. In *The Adaptive Web*, pages 377–408. Springer Berlin Heidelberg.
- Cortizo, J. C., Carrero, F. M., and Monsalve, B. (2010). An architecture for a general purpose multi-algorithm recommender system. In *Workshop on the Practical Use of Recommender Systems, Algorithms and Technologies (PRSAT 2010)*, page 51. Citeseer.
- easyrec (2009-2014). <http://www.easyrec.org/>. Accessed April 07, 2014.
- Erl, T. (2008). *Soa: principles of service design*, volume 1. Prentice Hall Upper Saddle River.
- express (2014). <http://expressjs.com>. Accessed April 07, 2014.

- Extensible Markup Language (XML) (2014). <http://www.w3.org/XML/>. Accessed April 07, 2014.
- Fielding, R. T. (2000). *Architectural styles and the design of network-based software architectures*. PhD thesis, University of California.
- Git (2014). <http://git-scm.com/>. Accessed April 07, 2014.
- Go (2014). <http://golang.org/>. Accessed April 07, 2014.
- Google Chrome (2014). <http://www.google.com/chrome/>. Accessed April 07, 2014.
- Hahsler, M. (2011). recommenderlab: A framework for developing and testing recommendation algorithms. *Nov*.
- Herlocker, J. L., Konstan, J. A., Terveen, L. G., and Riedl, J. T. (2004). Evaluating collaborative filtering recommender systems. *ACM Transactions on Information Systems*, 22(1):5–53.
- HTTP Vocabulary in RDF 1.0 (2014). <http://www.w3.org/TR/HTTP-in-RDF10/>. Accessed April 07, 2014.
- JavaScript Object Notation (JSON) (2014). <http://www.json.org/>. Accessed April 07, 2014.
- Linden, G., Smith, B., and York, J. (2003). Amazon.com recommendations: Item-to-item collaborative filtering. *Internet Computing, IEEE*, 7(1):76–80.
- Manouselis, N. and Costopoulou, C. (2007). Analysis and classification of multi-criteria recommender systems. *World Wide Web*, 10(4):415–441.
- McCarthy, J. (2007). What is artificial intelligence? <http://www-formal.stanford.edu/jmc/whatisai/whatisai.html>. Accessed April 07, 2014.
- Neo4j (2014). <http://www.neo4j.org/>. Accessed April 07, 2014.
- NewRelic (2014). <http://newrelic.com/>. Accessed April 07, 2014.
- node.js (2014). <http://nodejs.org/>. Accessed April 07, 2014.
- Prudsys (2014). prudsys realtime decisioning engine. <http://www.prudsys.de/en/solutions/rde-recommendations.html>. Accessed April 07, 2014.

- RabbitMQ (2014). <https://www.rabbitmq.com/>. Accessed April 07, 2014.
- Rack, C., Arbanowski, S., and Steglich, S. (2007). A generic multipurpose recommender system for contextual recommendations. In *Autonomous Decentralized Systems, 2007. ISADS'07. Eighth International Symposium on*, pages 445–450. IEEE.
- Redis (2014). <http://redis.io/>. Accessed April 07, 2014.
- Ricci, F., Rokach, L., and Shapira, B. (2011). *Introduction to Recommender Systems Handbook*. Springer US.
- Segaran, T. (2007). *Programming Collective Intelligence: Building Smart Web 2.0 Applications*. O'Reilly Media, Inc.
- The Apache Software Foundation (1999-2014). <http://www.apache.org/>. Accessed April 07, 2014.
- V8 JavaScript Engine (2014). <https://code.google.com/p/v8/>. Accessed April 07, 2014.
- VirtualBox (2014). <https://www.virtualbox.org/>. Accessed April 07, 2014.